

▼ Loading the data

▼ Imports and Google Drive mounting

```
try:  
    # %tensorflow_version only exists in Colab.  
    %tensorflow_version 2.x  
except Exception:  
    pass
```

```
# TensorFlow and tf.keras  
import tensorflow as tf  
from tensorflow import keras
```

```
print(tf.__version__)
```

```
# Helper libraries  
import numpy as np  
import matplotlib.pyplot as plt  
import matplotlib.gridspec as gridspec  
import sklearn as sk  
import pandas as pd  
from pandas import read_csv  
from datetime import datetime  
import math  
import os
```

```
# fix random seed for reproducibility  
seed = 2020  
np.random.seed(seed)
```

```

from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_absolute_error

from tensorflow.keras.datasets import cifar100
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout, Flatten, BatchNormalization, Activation
from tensorflow.keras.layers import Conv1D, MaxPooling1D, GlobalMaxPooling1D
from tensorflow.keras.layers import LSTM, GRU
from tensorflow.keras.regularizers import L1L2
from tensorflow.keras.constraints import max_norm
from tensorflow.keras.optimizers import Adam, SGD
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint
from tensorflow.keras.models import load_model
from tensorflow.keras import optimizers
from tensorflow.keras import Input
from tensorflow.keras.models import Model
from tensorflow.keras.constraints import MaxNorm

```

📄 TensorFlow 2.x selected.
2.1.0

```

from google.colab import drive
drive.mount('/content/gdrive')

!ls '/content/gdrive/My Drive/Colab Notebooks/DL2020/GA3/'

```

📄 Drive already mounted at /content/gdrive; to attempt to forcibly remount, call drive.mount("/content/gdrive", force_remount=True)

```

checkpoints                preprocessed_train_data_all.csv
pollution_data.csv         preprocessed_train_data.csv
preprocessed_pollution_data.csv preprocessed_val_data.csv
preprocessed_test_data.csv

```

▼ Loading the dataset

```
# this code only loads train and test datasets
```

```
# remember you also need to split off a validation dataset
# either by reusing and adapting the splitting code from the preprocessing notebook
# or by writing your own code
# like in the previous assignments, this means you should have 4 sets:

# train_all (train+validate), train, validate and test

# The PATH setting below assumes you just uploaded the data file to your Colab session
# When using Drive: replace this by the path where you put the data file
DATAPATH = '/content/gdrive/My Drive/Colab Notebooks/DL2020/GA3/'
TRAINDATAALLFILE = DATAPATH+'preprocessed_train_data_all.csv'
TRAINDATAFILE = DATAPATH+'preprocessed_train_data.csv'
VALDATAFILE = DATAPATH+'preprocessed_val_data.csv'
TESTDATAFILE = DATAPATH+'preprocessed_test_data.csv'

train_all_dataset = read_csv(TRAINDATAALLFILE, header=0, index_col=0)
train_dataset = read_csv(TRAINDATAFILE, header=0, index_col=0)
val_dataset = read_csv(VALDATAFILE, header=0, index_col=0)
test_dataset = read_csv(TESTDATAFILE, header=0, index_col=0)

#check - note that you can make many other plots directly from pandas dataframes:
# https://pandas.pydata.org/pandas-docs/version/0.16/visualization.html
train_dataset.loc[:, train_dataset.columns != 'Year'].plot(subplots=True,figsize=(16,26))
```




```

/usr/local/lib/python3.6/dist-packages/pandas/plotting/_matplotlib/tools.py:307: MatplotlibDeprecationWarning:
The rowNum attribute was deprecated in Matplotlib 3.2 and will be removed two minor releases later. Use ax.get_subplotspec().row
    layout[ax.rowNum, ax.colNum] = ax.get_visible()
/usr/local/lib/python3.6/dist-packages/pandas/plotting/_matplotlib/tools.py:307: MatplotlibDeprecationWarning:
The colNum attribute was deprecated in Matplotlib 3.2 and will be removed two minor releases later. Use ax.get_subplotspec().col
    layout[ax.rowNum, ax.colNum] = ax.get_visible()
/usr/local/lib/python3.6/dist-packages/pandas/plotting/_matplotlib/tools.py:313: MatplotlibDeprecationWarning:
The rowNum attribute was deprecated in Matplotlib 3.2 and will be removed two minor releases later. Use ax.get_subplotspec().row
    if not layout[ax.rowNum + 1, ax.colNum]:
/usr/local/lib/python3.6/dist-packages/pandas/plotting/_matplotlib/tools.py:313: MatplotlibDeprecationWarning:
The colNum attribute was deprecated in Matplotlib 3.2 and will be removed two minor releases later. Use ax.get_subplotspec().col
    if not layout[ax.rowNum + 1, ax.colNum]:
/usr/local/lib/python3.6/dist-packages/pandas/plotting/_matplotlib/tools.py:307: MatplotlibDeprecationWarning:
The rowNum attribute was deprecated in Matplotlib 3.2 and will be removed two minor releases later. Use ax.get_subplotspec().row
    layout[ax.rowNum, ax.colNum] = ax.get_visible()
/usr/local/lib/python3.6/dist-packages/pandas/plotting/_matplotlib/tools.py:307: MatplotlibDeprecationWarning:
The colNum attribute was deprecated in Matplotlib 3.2 and will be removed two minor releases later. Use ax.get_subplotspec().col
    layout[ax.rowNum, ax.colNum] = ax.get_visible()
/usr/local/lib/python3.6/dist-packages/pandas/plotting/_matplotlib/tools.py:313: MatplotlibDeprecationWarning:
The rowNum attribute was deprecated in Matplotlib 3.2 and will be removed two minor releases later. Use ax.get_subplotspec().row
    if not layout[ax.rowNum + 1, ax.colNum]:
/usr/local/lib/python3.6/dist-packages/pandas/plotting/_matplotlib/tools.py:313: MatplotlibDeprecationWarning:
The colNum attribute was deprecated in Matplotlib 3.2 and will be removed two minor releases later. Use ax.get_subplotspec().col
    if not layout[ax.rowNum + 1, ax.colNum]:
array([<matplotlib.axes._subplots.AxesSubplot object at 0x7f3290398438>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f329036d6d8>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f32903246d8>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f32902da6d8>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f329028f6d8>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f32902c36d8>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f32902796d8>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f329022e6a0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f329022e710>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f32901986d8>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f329014d6d8>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f32901816d8>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f32901366d8>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f32900eb6d8>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f32900a06d8>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f32900556d8>],
      dtype=object)

```





▼ Extract train, validate and test features and labels

```
# extract data from dataframes, selecting the features you want
```

```

# note that the labels simply contain the pollution data for now
# depending on the window sizes used in training, the correct values will be cut out
features = np.arange(1,17)
pollution = 1

# Important: Tensorflow 2.x gives an error if you omit
# the np.asarray(...,dtype=np.float32)
train_all_values = np.asarray(train_all_dataset.values[:,features],dtype=np.float32)
train_all_labels = np.asarray(train_all_dataset.values[:,pollution],dtype=np.float32)
train_values = np.asarray(train_dataset.values[:,features],dtype=np.float32)
train_labels = np.asarray(train_dataset.values[:,pollution],dtype=np.float32)
val_values = np.asarray(val_dataset.values[:,features],dtype=np.float32)
val_labels = np.asarray(val_dataset.values[:,pollution],dtype=np.float32)
test_values = np.asarray(test_dataset.values[:,features],dtype=np.float32)
test_labels = np.asarray(test_dataset.values[:,pollution],dtype=np.float32)

```

▼ Normalising the data

```

# imports to show that there are many different scalers
# especially with recurrent NNs, the choice of scaler can make a difference
# look up what they do before choosing which one to try
# from sklearn.preprocessing import StandardScaler
# from sklearn.preprocessing import MinMaxScaler
# from sklearn.preprocessing import MaxAbsScaler
# from sklearn.preprocessing import RobustScaler
# from sklearn.preprocessing import Normalizer # BAD for this case
# from sklearn.preprocessing import QuantileTransformer
from sklearn.preprocessing import PowerTransformer # BEST for this case

# Example: train standard scalers, apply to train and test data
# adapt to do all you need to do ...
SS1 = PowerTransformer()
SS1.fit(train_values)

train_scaled = SS1.transform(train_values)

```

```
val_scaled = SS1.transform(val_values)

SS2 = PowerTransformer()
SS2.fit(train_all_values)

train_all_scaled = SS2.transform(train_all_values)
test_scaled = SS2.transform(test_values)

pd.DataFrame(train_scaled, columns=train_dataset.columns[1:]).plot(subplots=True, figsize=(16,26))
```






```

/usr/local/lib/python3.6/dist-packages/numpy/core/_methods.py:195: RuntimeWarning: overflow encountered in multiply
  x = um.multiply(x, x, out=x)
/usr/local/lib/python3.6/dist-packages/numpy/core/_methods.py:195: RuntimeWarning: overflow encountered in multiply
  x = um.multiply(x, x, out=x)
/usr/local/lib/python3.6/dist-packages/pandas/plotting/_matplotlib/tools.py:307: MatplotlibDeprecationWarning:
The rowNum attribute was deprecated in Matplotlib 3.2 and will be removed two minor releases later. Use ax.get_subplotspec().row
  layout[ax.rowNum, ax.colNum] = ax.get_visible()
/usr/local/lib/python3.6/dist-packages/pandas/plotting/_matplotlib/tools.py:307: MatplotlibDeprecationWarning:
The colNum attribute was deprecated in Matplotlib 3.2 and will be removed two minor releases later. Use ax.get_subplotspec().col
  layout[ax.rowNum, ax.colNum] = ax.get_visible()
/usr/local/lib/python3.6/dist-packages/pandas/plotting/_matplotlib/tools.py:313: MatplotlibDeprecationWarning:
The rowNum attribute was deprecated in Matplotlib 3.2 and will be removed two minor releases later. Use ax.get_subplotspec().row
  if not layout[ax.rowNum + 1, ax.colNum]:
/usr/local/lib/python3.6/dist-packages/pandas/plotting/_matplotlib/tools.py:313: MatplotlibDeprecationWarning:
The colNum attribute was deprecated in Matplotlib 3.2 and will be removed two minor releases later. Use ax.get_subplotspec().col
  if not layout[ax.rowNum + 1, ax.colNum]:
/usr/local/lib/python3.6/dist-packages/pandas/plotting/_matplotlib/tools.py:307: MatplotlibDeprecationWarning:
The rowNum attribute was deprecated in Matplotlib 3.2 and will be removed two minor releases later. Use ax.get_subplotspec().row
  layout[ax.rowNum, ax.colNum] = ax.get_visible()
/usr/local/lib/python3.6/dist-packages/pandas/plotting/_matplotlib/tools.py:307: MatplotlibDeprecationWarning:
The colNum attribute was deprecated in Matplotlib 3.2 and will be removed two minor releases later. Use ax.get_subplotspec().col
  layout[ax.rowNum, ax.colNum] = ax.get_visible()
/usr/local/lib/python3.6/dist-packages/pandas/plotting/_matplotlib/tools.py:313: MatplotlibDeprecationWarning:
The rowNum attribute was deprecated in Matplotlib 3.2 and will be removed two minor releases later. Use ax.get_subplotspec().row
  if not layout[ax.rowNum + 1, ax.colNum]:
/usr/local/lib/python3.6/dist-packages/pandas/plotting/_matplotlib/tools.py:313: MatplotlibDeprecationWarning:
The colNum attribute was deprecated in Matplotlib 3.2 and will be removed two minor releases later. Use ax.get_subplotspec().col
  if not layout[ax.rowNum + 1, ax.colNum]:
array([<matplotlib.axes._subplots.AxesSubplot object at 0x7f328f961198>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f328f980f98>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f328e124f98>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f328e0d8eb8>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f328e098128>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f328e049358>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f32903249e8>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f32902eb080>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f32902eb198>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f328f9d5400>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f328e138978>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f328df5e400>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f328df13400>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f328dec9400>])

```

```
....._subplots.AxesSubplot object at 0x7f328de7f400>,\n<matplotlib.axes._subplots.AxesSubplot object at 0x7f328deb3400>],\ndtype=object)
```





▼ Dataset manipulations

```
# dense network:
# here you only need to take into account that we are predicting 6 steps ahead
# this means that the features of the first timestep (index 0)
# are used to predict the 7th pollution value (index 6) so the first 6 and the last 6 feature-samples labels are omitted
def create_dataset_dense(train, test, ahead=6): # can use this with different 'ahead' values, but default is set to 6
    return train[:-ahead,:], test[ahead:]

# window-based and recurrent networks:
# now, you use a window of k history values to predict
# this means that the features of the first k timesteps (indices 0 to k-1)
# are used to predict the k+6th pollution value (index k+6-1)
# output dimension of train data is samples x window x features
def create_dataset_windowed(train, test, ahead=6, window_size=1):
    samples = train.shape[0]-ahead-(window_size-1)
    dataX= []
    for i in range(samples):
        a = train[i:(i+window_size), :]
        dataX.append(a)

    return np.array(dataX), test[ahead+window_size-1:]

# reshape arrays so that they can be used as input for the top model of the ensemble
def create_dataset_top(train, test, ahead=6, window_size=1):
    return np.concatenate(train, axis=1), test[ahead+window_size-1:]
```

▼ Help functions for visualization

```
def plot_history(history):
    plt.figure(figsize=(6,4))

    plt.xlabel('Epoch')
    plt.ylabel('Mae')
```

```

plt.plot(history.epoch, np.array(history.history['mae']), 'g-',
         label='Train MAE')
try:
    plt.plot(history.epoch, np.array(history.history['val_mae']), 'r-',
             label='Validation MAE')
except:
    pass
plt.legend()
plt.show()

```

```

def PlotResults(labels, predictions, binsize=10):
    num_samples = len(labels)

    fig = plt.figure(figsize=(16,16))
    spec = gridspec.GridSpec(ncols=4, nrows=4, figure=fig)
    ax1 = fig.add_subplot(spec[0, :])
    ax2 = fig.add_subplot(spec[1, :])
    ax3 = fig.add_subplot(spec[2:,0:2])
    ax4 = fig.add_subplot(spec[2:,2:])

    ax1.plot(labels, 'k-', label='Labels')
    ax1.plot(predictions, 'r-', label='Predictions')
    ax1.set_ylabel('Pollution')
    ax1.legend()

    errors = np.absolute(labels - predictions)
    ax2.plot(errors, 'k-')
    ax2.set_ylabel("Absolute error")

    ax3.scatter(labels, predictions)
    ax3.set_xlabel('Labels')
    ax3.set_ylabel('Predictions')

    bins = np.arange(0, (np.ceil(np.max(errors)/binsize)+1)*binsize, binsize)

    ax4.hist(errors, bins=bins)
    ax4.set_xlabel('Absolute error')
    ax4.set_ylabel('Frequency')

```

```
ax4.set_ylabel('Frequency')
```

```
plt.show()
```

▼ Step 1

```
X_train_d, r_train_d = create_dataset_dense(train_scaled, train_labels)
X_val_d, r_val_d = create_dataset_dense(val_scaled, val_labels)
X_train_all_d, r_train_all_d = create_dataset_dense(train_all_scaled, train_all_labels)
X_test_d, r_test_d = create_dataset_dense(test_scaled, test_labels)

# Train a Dense network on the scaled features
def dense_model(learning_rate=0.01, hidden=[256,1024,1024,128], dropouts=[0.5,0.5,0.5,0.1]):
    # create linear model
    model = Sequential()
    model.add(Dense(hidden[0], kernel_initializer='he_uniform', input_shape=(train_scaled.shape[1],)))
    model.add(BatchNormalization())
    model.add(Activation('relu'))
    model.add(Dropout(dropouts[0]))

    for idx in range(1, len(hidden)):
        model.add(Dense(hidden[idx], kernel_initializer='he_uniform'))
        model.add(BatchNormalization())
        model.add(Activation('relu'))
        model.add(Dropout(dropouts[idx]))

    model.add(Dense(1))
    model.add(Activation('linear'))
    # no dropout at the end

    optim = tf.keras.optimizers.Adam(lr=learning_rate)
    model.compile(loss='mae',
                  optimizer=optim,
                  # keep extra metrics: mse and mae without regularisation terms
                  metrics=['mse', 'mae'])

    return model
```

```

model = dense_model()
model.summary()

batch_size = 5*168
epochs = 70

# cp_path = '/content/gdrive/My Drive/Colab Notebooks/DL2020/GA3/checkpoints/step_1/cp.ckpt'
# cp_callback = tf.keras.callbacks.ModelCheckpoint(cp_path,
#                                                  monitor='val_mae',
#                                                  verbose=1,
#                                                  save_best_only=True,
#                                                  save_weights_only=True)
# lr_callback = tf.keras.callbacks.ReduceLROnPlateau(monitor='val_mae',
#                                                    factor=0.9,
#                                                    patience=5,
#                                                    verbose=1,
#                                                    min_lr=0.0001)

# history = model.fit(X_train_d,
#                     r_train_d,
#                     shuffle=True,
#                     batch_size=batch_size,
#                     epochs=epochs,
#                     validation_data=(X_val_d,r_val_d),
#                     callbacks=[cp_callback, lr_callback])
# model.load_weights(cp_path)

history = model.fit(X_train_all_d,
                    r_train_all_d,
                    shuffle=True,
                    batch_size=batch_size,
                    epochs=epochs)

```



Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
dense (Dense)	(None, 256)	4352
batch_normalization (Batch Normalization)	(None, 256)	1024
activation (Activation)	(None, 256)	0
dropout (Dropout)	(None, 256)	0
dense_1 (Dense)	(None, 1024)	263168
batch_normalization_1 (Batch Normalization)	(None, 1024)	4096
activation_1 (Activation)	(None, 1024)	0
dropout_1 (Dropout)	(None, 1024)	0
dense_2 (Dense)	(None, 1024)	1049600
batch_normalization_2 (Batch Normalization)	(None, 1024)	4096
activation_2 (Activation)	(None, 1024)	0
dropout_2 (Dropout)	(None, 1024)	0
dense_3 (Dense)	(None, 128)	131200
batch_normalization_3 (Batch Normalization)	(None, 128)	512
activation_3 (Activation)	(None, 128)	0
dropout_3 (Dropout)	(None, 128)	0
dense_4 (Dense)	(None, 1)	129
activation_4 (Activation)	(None, 1)	0
=====		
Total params: 1,458,177		
Trainable params: 1,453,313		

Non-trainable params: 4,864

Train on 35034 samples

Epoch 1/70

35034/35034 [=====] - 2s 46us/sample - loss: 78.5146 - mse: 13661.8945 - mae: 78.5146

Epoch 2/70

35034/35034 [=====] - 0s 9us/sample - loss: 44.0658 - mse: 5203.6519 - mae: 44.0658

Epoch 3/70

35034/35034 [=====] - 0s 8us/sample - loss: 37.3092 - mse: 3571.8965 - mae: 37.3092

Epoch 4/70

35034/35034 [=====] - 0s 8us/sample - loss: 36.7322 - mse: 3503.4966 - mae: 36.7322

Epoch 5/70

35034/35034 [=====] - 0s 8us/sample - loss: 36.3990 - mse: 3457.7317 - mae: 36.3990

Epoch 6/70

35034/35034 [=====] - 0s 8us/sample - loss: 36.0665 - mse: 3418.4182 - mae: 36.0665

Epoch 7/70

35034/35034 [=====] - 0s 8us/sample - loss: 35.9245 - mse: 3398.2158 - mae: 35.9245

Epoch 8/70

35034/35034 [=====] - 0s 8us/sample - loss: 35.8215 - mse: 3389.9365 - mae: 35.8215

Epoch 9/70

35034/35034 [=====] - 0s 8us/sample - loss: 35.6267 - mse: 3354.0217 - mae: 35.6267

Epoch 10/70

35034/35034 [=====] - 0s 8us/sample - loss: 35.4905 - mse: 3326.1458 - mae: 35.4905

Epoch 11/70

35034/35034 [=====] - 0s 8us/sample - loss: 35.2678 - mse: 3320.1406 - mae: 35.2678

Epoch 12/70

35034/35034 [=====] - 0s 8us/sample - loss: 35.1846 - mse: 3303.9519 - mae: 35.1846

Epoch 13/70

35034/35034 [=====] - 0s 8us/sample - loss: 35.2897 - mse: 3323.4863 - mae: 35.2897

Epoch 14/70

35034/35034 [=====] - 0s 8us/sample - loss: 35.1027 - mse: 3303.5291 - mae: 35.1027

Epoch 15/70

35034/35034 [=====] - 0s 9us/sample - loss: 34.8917 - mse: 3261.8740 - mae: 34.8917

Epoch 16/70

35034/35034 [=====] - 0s 8us/sample - loss: 34.8479 - mse: 3262.4751 - mae: 34.8479

Epoch 17/70

35034/35034 [=====] - 0s 8us/sample - loss: 34.6774 - mse: 3225.4436 - mae: 34.6774

Epoch 18/70

35034/35034 [=====] - 0s 8us/sample - loss: 34.6194 - mse: 3225.2275 - mae: 34.6194

Epoch 19/70

35034/35034 [=====] - 0s 8us/sample - loss: 34.7590 - mse: 3252.2400 - mae: 34.7590

Epoch 20/70

35034/35034 [=====] - 0s 8us/sample - loss: 34.5515 - mse: 3195.1504 - mae: 34.5515
Epoch 21/70
35034/35034 [=====] - 0s 9us/sample - loss: 34.4425 - mse: 3205.3689 - mae: 34.4425
Epoch 22/70
35034/35034 [=====] - 0s 9us/sample - loss: 34.3655 - mse: 3200.5151 - mae: 34.3655
Epoch 23/70
35034/35034 [=====] - 0s 8us/sample - loss: 34.3033 - mse: 3168.3855 - mae: 34.3033
Epoch 24/70
35034/35034 [=====] - 0s 9us/sample - loss: 34.3860 - mse: 3185.9153 - mae: 34.3860
Epoch 25/70
35034/35034 [=====] - 0s 9us/sample - loss: 34.2461 - mse: 3172.3406 - mae: 34.2461
Epoch 26/70
35034/35034 [=====] - 0s 8us/sample - loss: 34.1995 - mse: 3165.9343 - mae: 34.1995
Epoch 27/70
35034/35034 [=====] - 0s 9us/sample - loss: 34.2764 - mse: 3162.7942 - mae: 34.2764
Epoch 28/70
35034/35034 [=====] - 0s 8us/sample - loss: 34.1207 - mse: 3155.5232 - mae: 34.1207
Epoch 29/70
35034/35034 [=====] - 0s 8us/sample - loss: 34.2727 - mse: 3143.4817 - mae: 34.2727
Epoch 30/70
35034/35034 [=====] - 0s 8us/sample - loss: 33.9865 - mse: 3148.1396 - mae: 33.9865
Epoch 31/70
35034/35034 [=====] - 0s 8us/sample - loss: 33.9636 - mse: 3133.0679 - mae: 33.9636
Epoch 32/70
35034/35034 [=====] - 0s 8us/sample - loss: 34.1207 - mse: 3154.8799 - mae: 34.1207
Epoch 33/70
35034/35034 [=====] - 0s 8us/sample - loss: 33.8180 - mse: 3113.9597 - mae: 33.8180
Epoch 34/70
35034/35034 [=====] - 0s 9us/sample - loss: 33.7960 - mse: 3107.2698 - mae: 33.7960
Epoch 35/70
35034/35034 [=====] - 0s 9us/sample - loss: 33.7567 - mse: 3114.9546 - mae: 33.7567
Epoch 36/70
35034/35034 [=====] - 0s 9us/sample - loss: 33.7672 - mse: 3091.2883 - mae: 33.7672
Epoch 37/70
35034/35034 [=====] - 0s 8us/sample - loss: 33.6124 - mse: 3075.5461 - mae: 33.6124
Epoch 38/70
35034/35034 [=====] - 0s 9us/sample - loss: 33.6878 - mse: 3084.3408 - mae: 33.6878
Epoch 39/70
35034/35034 [=====] - 0s 9us/sample - loss: 33.6742 - mse: 3105.8569 - mae: 33.6742
Epoch 40/70
35034/35034 [=====] - 0s 8us/sample - loss: 33.7394 - mse: 3085.9954 - mae: 33.7394
Epoch 41/70

35034/35034 [=====] - 0s 8us/sample - loss: 33.5806 - mse: 3050.9263 - mae: 33.5806
Epoch 42/70
35034/35034 [=====] - 0s 9us/sample - loss: 33.3304 - mse: 3024.0122 - mae: 33.3304
Epoch 43/70
35034/35034 [=====] - 0s 9us/sample - loss: 33.4608 - mse: 3065.6443 - mae: 33.4608
Epoch 44/70
35034/35034 [=====] - 0s 8us/sample - loss: 33.5393 - mse: 3051.9995 - mae: 33.5393
Epoch 45/70
35034/35034 [=====] - 0s 8us/sample - loss: 33.3328 - mse: 3041.8708 - mae: 33.3328
Epoch 46/70
35034/35034 [=====] - 0s 8us/sample - loss: 33.3387 - mse: 3034.3071 - mae: 33.3387
Epoch 47/70
35034/35034 [=====] - 0s 8us/sample - loss: 33.3511 - mse: 3040.2632 - mae: 33.3511
Epoch 48/70
35034/35034 [=====] - 0s 8us/sample - loss: 33.3621 - mse: 3051.0061 - mae: 33.3621
Epoch 49/70
35034/35034 [=====] - 0s 9us/sample - loss: 33.2823 - mse: 3026.6216 - mae: 33.2823
Epoch 50/70
35034/35034 [=====] - 0s 8us/sample - loss: 33.3578 - mse: 3058.5618 - mae: 33.3578
Epoch 51/70
35034/35034 [=====] - 0s 8us/sample - loss: 33.2197 - mse: 3014.4700 - mae: 33.2197
Epoch 52/70
35034/35034 [=====] - 0s 8us/sample - loss: 33.3109 - mse: 3045.2559 - mae: 33.3109
Epoch 53/70
35034/35034 [=====] - 0s 8us/sample - loss: 33.2751 - mse: 3025.4312 - mae: 33.2751
Epoch 54/70
35034/35034 [=====] - 0s 8us/sample - loss: 33.2664 - mse: 3006.1978 - mae: 33.2664
Epoch 55/70
35034/35034 [=====] - 0s 8us/sample - loss: 33.1946 - mse: 3017.5742 - mae: 33.1946
Epoch 56/70
35034/35034 [=====] - 0s 8us/sample - loss: 33.0301 - mse: 2993.9436 - mae: 33.0301
Epoch 57/70
35034/35034 [=====] - 0s 9us/sample - loss: 33.0644 - mse: 2972.7712 - mae: 33.0644
Epoch 58/70
35034/35034 [=====] - 0s 8us/sample - loss: 33.1257 - mse: 2992.3965 - mae: 33.1257
Epoch 59/70
35034/35034 [=====] - 0s 8us/sample - loss: 33.1044 - mse: 3015.5137 - mae: 33.1044
Epoch 60/70
35034/35034 [=====] - 0s 9us/sample - loss: 33.0605 - mse: 2999.5271 - mae: 33.0605
Epoch 61/70
35034/35034 [=====] - 0s 8us/sample - loss: 32.9215 - mse: 2969.3818 - mae: 32.9215
Epoch 62/70

```
35034/35034 [=====] - 0s 8us/sample - loss: 33.1054 - mse: 2992.5251 - mae: 33.1054
Epoch 63/70
35034/35034 [=====] - 0s 8us/sample - loss: 32.9936 - mse: 2974.5037 - mae: 32.9936
Epoch 64/70
35034/35034 [=====] - 0s 9us/sample - loss: 32.8913 - mse: 2953.0444 - mae: 32.8913
Epoch 65/70
35034/35034 [=====] - 0s 8us/sample - loss: 32.8328 - mse: 2951.7742 - mae: 32.8328
Epoch 66/70
35034/35034 [=====] - 0s 8us/sample - loss: 32.9495 - mse: 2997.6707 - mae: 32.9495
Epoch 67/70
35034/35034 [=====] - 0s 9us/sample - loss: 32.8606 - mse: 2957.0000 - mae: 32.8606
Epoch 68/70
35034/35034 [=====] - 0s 9us/sample - loss: 32.9068 - mse: 2975.1118 - mae: 32.9068
Epoch 69/70
35034/35034 [=====] - 0s 8us/sample - loss: 32.7313 - mse: 2946.4006 - mae: 32.7313
Epoch 70/70
35034/35034 [=====] - 0s 8us/sample - loss: 32.6970 - mse: 2945.5493 - mae: 32.6970
```

```
plot_history(history)

# y_train = model.predict(X_train_d)
# y_val = model.predict(X_val_d)
y_train_all = model.predict(X_train_all_d)
y_test = model.predict(X_test_d)


# mae_train = mean_absolute_error(r_train_d,y_train)
# mae_val = mean_absolute_error(r_val_d,y_val)
mae_train_all = mean_absolute_error(r_train_all_d,y_train_all)
mae_test = mean_absolute_error(r_test_d,y_test)

# print(f"train      mae: {mae_train_all}")
# print(f"test mae: {mae_test}")
print(f"train      mae: {mae_train_all}")
print(f"validation mae: {mae_test}")



# Visualise first 1000 predictions for validation and test
# PlotResults(r_val_d[:1000],y_val[:1000,0])
```


```
PlotResults(r_test_d[:1000],y_test[:1000,0])
```





train mae: 30.360923767089844
validation mae: 34.866722106933594





▼ LIME

```
# predictions = model.predict(X_val_d)
# r_val_d_T = np.transpose(np.array([r_val_d]))
# mae_val = abs(r_val_d_T-predictions)
# mae_val_partitioned = np.argpartition(mae_val,-5,axis=0)
# largest = mae_val_partitioned[-5:]
# smallest = mae_val_partitioned[:5]
```

```
# print(f'Largest mistakes: {mae_val[largest]}')
# print(f'Smallest mistakes: {mae_val[smallest]}')

predictions = model.predict(X_test_d)
r_test_d_T = np.transpose(np.array([r_test_d]))
mae_test = abs(r_test_d_T-predictions)
mae_test_partitioned = np.argpartition(mae_test,-5,axis=0)
largest = mae_test_partitioned[-5:]
smallest = mae_test_partitioned[:5]
print(f'Largest mistakes: {mae_test[largest]}')
print(f'Smallest mistakes: {mae_test[smallest]}')
```

```
↳ Largest mistakes: [[419.20596]]

[[451.2628 ]]

[[516.565  ]]

[[437.78207]]

[[460.53723]]
Smallest mistakes: [[[ 4.7851257]]

[[ 3.4017372]]

[[ 8.788279  ]]

[[12.0851345]]

[[12.711214  ]]]
```

```
!pip install lime

import lime
import lime.lime_tabular

feature_names = ['Pollution',
                 'seasons_sin',
                 'seasons_cos',
```



```

'hour_of_week_sin',
'hour_of_week_cos',
'hour_of_day_sin',
'hour_of_day_cos',
'NE_strength',
'NW_strength',
'SE_strength',
'cv_strength',
'dew',
'temp',
'press',
'snow',
'rain']

```

```
explainer = lime.lime_tabular.LimeTabularExplainer(train_scaled, feature_names=feature_names, class_names=['Pollution'], verbose=True)
```

```

↳ Requirement already satisfied: lime in /usr/local/lib/python3.6/dist-packages (0.1.1.37)
Requirement already satisfied: scikit-image>=0.12; python_version >= "3.0" in /usr/local/lib/python3.6/dist-packages (from lime)
Requirement already satisfied: scipy in /usr/local/lib/python3.6/dist-packages (from lime) (1.4.1)
Requirement already satisfied: numpy in /usr/local/lib/python3.6/dist-packages (from lime) (1.18.2)
Requirement already satisfied: matplotlib; python_version >= "3.0" in /usr/local/lib/python3.6/dist-packages (from lime) (3.2.0)
Requirement already satisfied: scikit-learn>=0.18 in /usr/local/lib/python3.6/dist-packages (from lime) (0.22.2.post1)
Requirement already satisfied: progressbar in /usr/local/lib/python3.6/dist-packages (from lime) (2.5)
Requirement already satisfied: networkx>=2.0 in /usr/local/lib/python3.6/dist-packages (from scikit-image>=0.12; python_version
Requirement already satisfied: pillow>=4.3.0 in /usr/local/lib/python3.6/dist-packages (from scikit-image>=0.12; python_version
Requirement already satisfied: PyWavelets>=0.4.0 in /usr/local/lib/python3.6/dist-packages (from scikit-image>=0.12; python_vers
Requirement already satisfied: imageio>=2.3.0 in /usr/local/lib/python3.6/dist-packages (from scikit-image>=0.12; python_version
Requirement already satisfied: cycycler>=0.10 in /usr/local/lib/python3.6/dist-packages (from matplotlib; python_version >= "3.0"-
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /usr/local/lib/python3.6/dist-packages (from matplotlib
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.6/dist-packages (from matplotlib; python_version >= "
Requirement already satisfied: python-dateutil>=2.1 in /usr/local/lib/python3.6/dist-packages (from matplotlib; python_version >
Requirement already satisfied: joblib>=0.11 in /usr/local/lib/python3.6/dist-packages (from scikit-learn>=0.18->lime) (0.14.1)
Requirement already satisfied: decorator>=4.3.0 in /usr/local/lib/python3.6/dist-packages (from networkx>=2.0->scikit-image>=0.1
Requirement already satisfied: six in /usr/local/lib/python3.6/dist-packages (from cycycler>=0.10->matplotlib; python_version >= "
Requirement already satisfied: setuptools in /usr/local/lib/python3.6/dist-packages (from kiwisolver>=1.0.1->matplotlib; python_

```

```

# Explain largest and smallest mistakes
print('Largest mistakes:')
for i in largest:

```


```
print(f'Correct pollution: {r_val_d[i]}')
exp = explainer.explain_instance(X_val_d[i[0]], model.predict, num_features=5, num_samples=5000)
exp.show_in_notebook(show_table=False)
exp.score

print('Smallest mistakes:')
for i in smallest:
    print(f'Correct pollution: {r_val_d[i]}')
    exp = explainer.explain_instance(X_val_d[i[0]], model.predict, num_features=5, num_samples=5000)
    exp.show_in_notebook(show_table=False)
    exp.score
```




Largest mistakes:
 Correct pollution: [289.]
 Intercept -9.073544251839309
 Prediction_local [91.86833571]
 Right: 266.85614

Predicted value




negative

positive




Correct pollution: [227.]
 Intercept 10.458152778999157
 Prediction_local [111.73522731]
 Right: 164.46582

Predicted value




negative

positive




Correct pollution: [14.]
 Intercept 56.26142526155941
 Prediction_local [-15.80824271]
 Right: 18.795448

Predicted value




negative

positive




Correct pollution: [201.]
 Intercept 9.095140395595095
 Prediction_local [87.39353018]
 Right: 177.33276

Predicted value




negative

positive




Correct pollution: [370.]
 Intercept -9.622402032133586
 Prediction_local [99.93343762]
 Right: 290.445

Predicted value




negative

positive




Smallest mistakes:
 Correct pollution: [59.]
 Intercept 28.656407253076047
 Prediction_local [30.40274477]
 Right: 64.03888

Predicted value




negative


positive




Correct pollution: [23.]
 Intercept 49.04832864562894
 Prediction_local [-16.13255713]
 Right: 59.256817

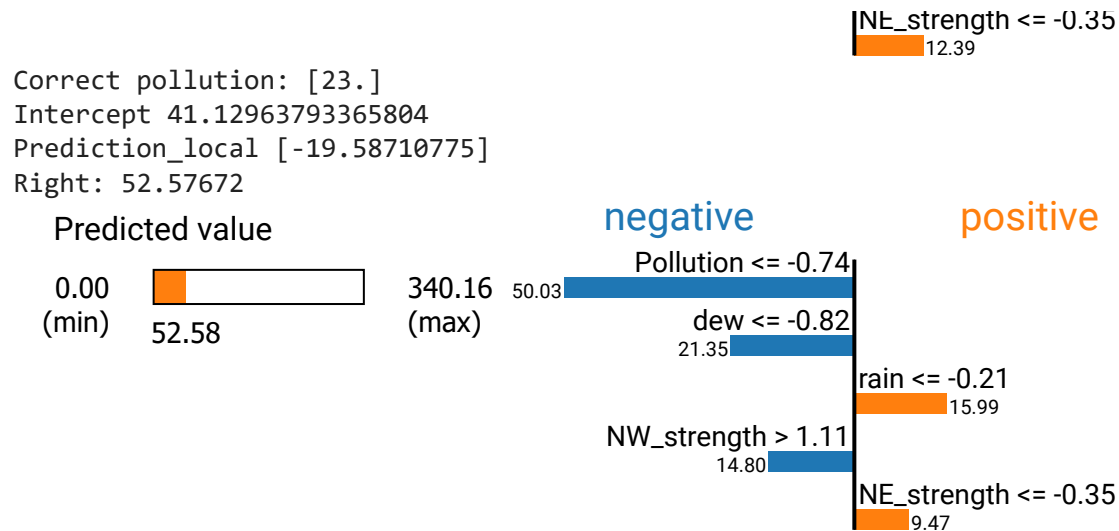


Correct pollution: [20.]
 Intercept 56.931562378563356
 Prediction_local [-16.65907207]
 Right: 50.335754



Correct pollution: [14.]
 Intercept 48.11350116855321
 Prediction_local [-3.73077516]
 Right: 24.017387





▼ Step 2

▼ Window-based CNN

```
# current code for pollution-only, add other features for multi-feature model,
# remove indices as in commented code if you want to use all features
train_features = [0,]

CNN_WINDOW_SIZE = 168
ahead = 6

# X_train_w,r_train_w = create_dataset_windowed(train_scaled[:,train_features], train_labels, window_size=CNN_WINDOW_SIZE)
# X_val_w,r_val_w = create_dataset_windowed(val_scaled[:,train_features], val_labels, window_size=CNN_WINDOW_SIZE)
X_train_all_w,r_train_all_w = create_dataset_windowed(train_all_scaled[:,train_features],train_all_labels, window_size=CNN_WINDOW_SIZE)
X_test_w,r_test_w = create_dataset_windowed(test_scaled[:,train_features],test_labels, window_size=CNN_WINDOW_SIZE)

def conv_model(CNN_WINDOW_SIZE=24, learning_rate=0.01, hidden=[64,128,256]):
    model = Sequential()
```

```

model.add(Conv1D(filters=hidden[0],
                 kernel_size=2,
                 padding='same',
                 kernel_initializer='he_uniform',
                 input_shape=(CNN_WINDOW_SIZE, len(train_features))))
model.add(BatchNormalization())
model.add(Activation('relu'))
model.add(MaxPooling1D(pool_size=2))

for idx in range(1, len(hidden)):
    model.add(Conv1D(filters=hidden[idx],
                     kernel_size=2,
                     padding='same',
                     kernel_initializer='he_uniform'))
    model.add(BatchNormalization())
    model.add(Activation('relu'))
    model.add(MaxPooling1D(pool_size=2))

model.add(Flatten())
model.add(Dense(1, kernel_initializer='glorot_uniform'))
model.add(Activation('linear'))

optim = tf.keras.optimizers.Adam(lr=learning_rate)
model.compile(loss='mae',
              optimizer=optim,
              # keep extra metrics: mse and mae without regularisation terms
              metrics=['mse', 'mae'])

return model

```

```

batch_size = 8*168 # 8 weeks
epochs = 30
learning_rate = 0.001

```

```

model = conv_model(CNN_WINDOW_SIZE=CNN_WINDOW_SIZE, learning_rate=learning_rate, hidden=[64,128,256])
model.summary()

```

```
# cp_path = '/content/gdrive/My Drive/Colab Notebooks/DL2020/GA3/checkpoints/step_2/cp.ckpt'
# cp_callback = tf.keras.callbacks.ModelCheckpoint(cp_path,
#
#                                     monitor='val_mae',
#                                     verbose=1,
#                                     save_best_only=True,
#                                     save_weights_only=True)

# history = model.fit(X_train_w, r_train_w,
#                     shuffle=True,
#                     batch_size=batch_size,
#                     epochs=epochs,
#                     validation_data=(X_val_w, r_val_w),
#                     callbacks=[cp_callback])
# model.load_weights(cp_path)

history = model.fit(X_train_all_w,
                    r_train_all_w,
                    shuffle=True,
                    batch_size=batch_size,
                    epochs=epochs)
```



Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
conv1d (Conv1D)	(None, 168, 64)	192
batch_normalization_4 (Batch Normalization)	(None, 168, 64)	256
activation_5 (Activation)	(None, 168, 64)	0
max_pooling1d (MaxPooling1D)	(None, 84, 64)	0
conv1d_1 (Conv1D)	(None, 84, 128)	16512
batch_normalization_5 (Batch Normalization)	(None, 84, 128)	512
activation_6 (Activation)	(None, 84, 128)	0
max_pooling1d_1 (MaxPooling1D)	(None, 42, 128)	0
conv1d_2 (Conv1D)	(None, 42, 256)	65792
batch_normalization_6 (Batch Normalization)	(None, 42, 256)	1024
activation_7 (Activation)	(None, 42, 256)	0
max_pooling1d_2 (MaxPooling1D)	(None, 21, 256)	0
flatten (Flatten)	(None, 5376)	0
dense_5 (Dense)	(None, 1)	5377
activation_8 (Activation)	(None, 1)	0

=====

Total params: 89,665
Trainable params: 88,769
Non-trainable params: 896

Train on 34867 samples

Epoch 1/30

34867/34867 [=====] - 5s 141us/sample - loss: 71.9665 - mse: 11722.4336 - mae: 71.9666

Epoch 2/30

34867/34867 [=====] - 2s 49us/sample - loss: 55.0473 - mse: 6389.7373 - mae: 55.0473
Epoch 3/30
34867/34867 [=====] - 2s 50us/sample - loss: 43.6378 - mse: 4429.9492 - mae: 43.6378
Epoch 4/30
34867/34867 [=====] - 2s 50us/sample - loss: 39.2266 - mse: 3833.0310 - mae: 39.2266
Epoch 5/30
34867/34867 [=====] - 2s 50us/sample - loss: 38.6068 - mse: 3780.6689 - mae: 38.6068
Epoch 6/30
34867/34867 [=====] - 2s 50us/sample - loss: 38.3187 - mse: 3751.9556 - mae: 38.3187
Epoch 7/30
34867/34867 [=====] - 2s 50us/sample - loss: 38.0931 - mse: 3731.8708 - mae: 38.0931
Epoch 8/30
34867/34867 [=====] - 2s 50us/sample - loss: 38.0126 - mse: 3713.8384 - mae: 38.0126
Epoch 9/30
34867/34867 [=====] - 2s 50us/sample - loss: 37.8579 - mse: 3701.4709 - mae: 37.8579
Epoch 10/30
34867/34867 [=====] - 2s 51us/sample - loss: 37.7775 - mse: 3684.1226 - mae: 37.7775
Epoch 11/30
34867/34867 [=====] - 2s 50us/sample - loss: 37.5990 - mse: 3665.1528 - mae: 37.5990
Epoch 12/30
34867/34867 [=====] - 2s 50us/sample - loss: 37.4350 - mse: 3645.4639 - mae: 37.4350
Epoch 13/30
34867/34867 [=====] - 2s 50us/sample - loss: 37.4208 - mse: 3638.4688 - mae: 37.4208
Epoch 14/30
34867/34867 [=====] - 2s 50us/sample - loss: 37.3966 - mse: 3634.8154 - mae: 37.3966
Epoch 15/30
34867/34867 [=====] - 2s 50us/sample - loss: 37.2579 - mse: 3627.5024 - mae: 37.2579
Epoch 16/30
34867/34867 [=====] - 2s 50us/sample - loss: 37.2016 - mse: 3617.1348 - mae: 37.2016
Epoch 17/30
34867/34867 [=====] - 2s 51us/sample - loss: 37.1703 - mse: 3611.6221 - mae: 37.1703
Epoch 18/30
34867/34867 [=====] - 2s 51us/sample - loss: 37.0839 - mse: 3592.8589 - mae: 37.0839
Epoch 19/30
34867/34867 [=====] - 2s 50us/sample - loss: 36.9892 - mse: 3588.2493 - mae: 36.9892
Epoch 20/30
34867/34867 [=====] - 2s 50us/sample - loss: 37.0456 - mse: 3583.0808 - mae: 37.0456
Epoch 21/30
34867/34867 [=====] - 2s 50us/sample - loss: 36.9462 - mse: 3572.5754 - mae: 36.9462
Epoch 22/30
34867/34867 [=====] - 2s 50us/sample - loss: 36.8073 - mse: 3556.1616 - mae: 36.8073
Epoch 23/30

```
34867/34867 [=====] - 2s 50us/sample - loss: 36.7746 - mse: 3558.0537 - mae: 36.7746
Epoch 24/30
34867/34867 [=====] - 2s 50us/sample - loss: 36.7149 - mse: 3546.3904 - mae: 36.7149
Epoch 25/30
34867/34867 [=====] - 2s 50us/sample - loss: 36.7548 - mse: 3548.9934 - mae: 36.7548
Epoch 26/30
34867/34867 [=====] - 2s 50us/sample - loss: 36.6818 - mse: 3527.5796 - mae: 36.6818
Epoch 27/30
34867/34867 [=====] - 2s 50us/sample - loss: 36.5393 - mse: 3526.1572 - mae: 36.5393
Epoch 28/30
34867/34867 [=====] - 2s 49us/sample - loss: 36.5183 - mse: 3508.5530 - mae: 36.5183
Epoch 29/30
34867/34867 [=====] - 2s 49us/sample - loss: 36.4431 - mse: 3505.9695 - mae: 36.4431
Epoch 30/30
34867/34867 [=====] - 2s 50us/sample - loss: 36.4045 - mse: 3503.3608 - mae: 36.4045
```


```
plot_history(history)

# y_train = model.predict(X_train_w)
# y_val = model.predict(X_val_w)
y_train_all = model.predict(X_train_all_w)
y_test = model.predict(X_test_w)

# mae_train = mean_absolute_error(r_train_w,y_train)
# mae_val = mean_absolute_error(r_val_w,y_val)
mae_train_all = mean_absolute_error(r_train_all_w,y_train_all)
mae_test = mean_absolute_error(r_test_w,y_test)



# print(f"train      mae: {mae_train}")
# print(f"validation mae: {mae_val}")
print(f"train all mae: {mae_train_all}")
print(f"test       mae: {mae_test}")


# Visualise first 1000 predictions for validation and test
# PlotResults(r_val_w[:1000],y_val[:1000,0])
PlotResults(r_test_w[:1000],y_test[:1000,0])
```



train all mae: 36.263267517089844

test mae: 36.93035888671875





▼ LSTM/GRU

```
# current code for pollution-only, add other features for multi-feature model,  
# remove indices as in commented code if you want to use all features  
train_features = [0,]
```

```
GRU_WINDOW_SIZE = 20  
ahead = 6
```

```

# For GRU, input has to be between 0 and 1
# max_value_train = np.max(train_scaled[:,train_features])
# min_value_train = np.min(train_scaled[:,train_features])
# max_value_val = np.max(val_scaled[:,train_features])
# min_value_val = np.min(val_scaled[:,train_features])
max_value_train_all = np.max(train_all_scaled[:,train_features])
min_value_train_all = np.min(train_all_scaled[:,train_features])
max_value_test = np.max(test_scaled[:,train_features])
min_value_test = np.min(test_scaled[:,train_features])

# train_scaled_gru = (train_scaled[:,train_features]-min_value_train)/(max_value_train-min_value_train)
# val_scaled_gru = (val_scaled[:,train_features]-min_value_val)/(max_value_val-min_value_val)
train_all_scaled_gru = (train_all_scaled[:,train_features]-min_value_train_all)/(max_value_train_all-min_value_train_all)
test_scaled_gru = (test_scaled[:,train_features]-min_value_test)/(max_value_test-min_value_test)

# X_train_gru,r_train_gru = create_dataset_windowed(train_scaled_gru, train_labels, window_size=GRU_WINDOW_SIZE)
# X_val_gru,r_val_gru = create_dataset_windowed(val_scaled_gru, val_labels, window_size=GRU_WINDOW_SIZE)
X_train_all_gru,r_train_all_gru = create_dataset_windowed(train_all_scaled_gru, train_all_labels, window_size=GRU_WINDOW_SIZE)
X_test_gru,r_test_gru = create_dataset_windowed(test_scaled_gru, test_labels, window_size=GRU_WINDOW_SIZE)

def gru_model(GRU_WINDOW_SIZE=24, units=512, learning_rate=0.001):
    model = Sequential()

    model.add(GRU(units=units,
                  return_sequences=False,
                  implementation=1,
                  input_shape=(GRU_WINDOW_SIZE,1)))
    model.add(Dropout(0.5))
    model.add(Dense(1))
    model.add(Activation('linear'))

    optim = tf.keras.optimizers.SGD(learning_rate=learning_rate,momentum=0.9,clipnorm=9)
    model.compile(loss='mae',
                  optimizer=optim,
                  # keep extra metrics: mse and mae without regularisation terms
                  metrics=['mse', 'mae'])

    return model

```

```

batch_size = 2*168
epochs = 50
learning_rate = 0.01

model = gru_model(GRU_WINDOW_SIZE=GRU_WINDOW_SIZE, units=512, learning_rate=learning_rate)
model.summary()

# cp_path = '/content/gdrive/My Drive/Colab Notebooks/DL2020/GA3/checkpoints/step_2/cp.ckpt'
# cp_callback = tf.keras.callbacks.ModelCheckpoint(cp_path,
#                                                  monitor='val_mae',
#                                                  verbose=1,
#                                                  save_best_only=True,
#                                                  save_weights_only=True)
# es_callback = tf.keras.callbacks.EarlyStopping(monitor='val_mae',
#                                                  min_delta=0.0001,
#                                                  patience=5)
# history = model.fit(X_train_gru,
#                     r_train_gru,
#                     shuffle=True,
#                     batch_size=batch_size,
#                     epochs=epochs,
#                     validation_data=(X_val_gru,r_val_gru),
#                     callbacks=[cp_callback, es_callback])
# model.load_weights(cp_path)

history = model.fit(X_train_all_gru,
                    r_train_all_gru,
                    shuffle=True,
                    batch_size=batch_size,
                    epochs=epochs)

```



Model: "sequential_2"

Layer (type)	Output Shape	Param #
gru (GRU)	(None, 512)	791040
dropout_4 (Dropout)	(None, 512)	0
dense_6 (Dense)	(None, 1)	513
activation_9 (Activation)	(None, 1)	0

Total params: 791,553

Trainable params: 791,553

Non-trainable params: 0

Train on 35015 samples

Epoch 1/50

35015/35015 [=====] - 2s 69us/sample - loss: 71.2187 - mse: 10740.4189 - mae: 71.2187

Epoch 2/50

35015/35015 [=====] - 1s 42us/sample - loss: 65.6855 - mse: 9135.8750 - mae: 65.6854

Epoch 3/50

35015/35015 [=====] - 1s 42us/sample - loss: 64.9366 - mse: 8905.2334 - mae: 64.9367

Epoch 4/50

35015/35015 [=====] - 1s 42us/sample - loss: 60.5093 - mse: 8061.3975 - mae: 60.5093

Epoch 5/50

35015/35015 [=====] - 1s 42us/sample - loss: 46.6006 - mse: 5129.3594 - mae: 46.6006

Epoch 6/50

35015/35015 [=====] - 1s 42us/sample - loss: 43.4271 - mse: 4405.7178 - mae: 43.4271

Epoch 7/50

35015/35015 [=====] - 1s 42us/sample - loss: 42.7156 - mse: 4272.8228 - mae: 42.7156

Epoch 8/50

35015/35015 [=====] - 2s 43us/sample - loss: 42.4921 - mse: 4234.6924 - mae: 42.4921

Epoch 9/50

35015/35015 [=====] - 1s 42us/sample - loss: 42.1491 - mse: 4189.7827 - mae: 42.1491

Epoch 10/50

35015/35015 [=====] - 2s 43us/sample - loss: 42.0081 - mse: 4172.7471 - mae: 42.0081

Epoch 11/50

35015/35015 [=====] - 1s 43us/sample - loss: 42.0503 - mse: 4174.9243 - mae: 42.0503

Epoch 12/50

35015/35015 [=====] - 1s 43us/sample - loss: 41.6807 - mse: 4139.4043 - mae: 41.6807

Epoch 13/50

35015/35015 [=====] - 1s 43us/sample - loss: 41.6494 - mse: 4130.2139 - mae: 41.6494
Epoch 14/50
35015/35015 [=====] - 1s 43us/sample - loss: 41.5756 - mse: 4127.7944 - mae: 41.5756
Epoch 15/50
35015/35015 [=====] - 2s 44us/sample - loss: 41.4433 - mse: 4117.9678 - mae: 41.4433
Epoch 16/50
35015/35015 [=====] - 2s 43us/sample - loss: 41.3191 - mse: 4095.5193 - mae: 41.3191
Epoch 17/50
35015/35015 [=====] - 1s 43us/sample - loss: 41.2442 - mse: 4095.8853 - mae: 41.2442
Epoch 18/50
35015/35015 [=====] - 2s 43us/sample - loss: 41.1335 - mse: 4095.1890 - mae: 41.1335
Epoch 19/50
35015/35015 [=====] - 2s 43us/sample - loss: 41.0751 - mse: 4077.7515 - mae: 41.0751
Epoch 20/50
35015/35015 [=====] - 2s 43us/sample - loss: 41.1148 - mse: 4081.1523 - mae: 41.1148
Epoch 21/50
35015/35015 [=====] - 1s 43us/sample - loss: 41.0450 - mse: 4079.3882 - mae: 41.0450
Epoch 22/50
35015/35015 [=====] - 2s 43us/sample - loss: 40.9255 - mse: 4050.8711 - mae: 40.9255
Epoch 23/50
35015/35015 [=====] - 2s 43us/sample - loss: 40.7574 - mse: 4035.2146 - mae: 40.7574
Epoch 24/50
35015/35015 [=====] - 2s 43us/sample - loss: 40.7147 - mse: 4037.5176 - mae: 40.7147
Epoch 25/50
35015/35015 [=====] - 1s 42us/sample - loss: 40.8396 - mse: 4037.1360 - mae: 40.8396
Epoch 26/50
35015/35015 [=====] - 2s 43us/sample - loss: 40.5780 - mse: 4010.3914 - mae: 40.5780
Epoch 27/50
35015/35015 [=====] - 1s 43us/sample - loss: 40.5598 - mse: 4028.8118 - mae: 40.5598
Epoch 28/50
35015/35015 [=====] - 1s 43us/sample - loss: 40.5789 - mse: 4019.4075 - mae: 40.5789
Epoch 29/50
35015/35015 [=====] - 2s 43us/sample - loss: 40.4464 - mse: 4007.7205 - mae: 40.4464
Epoch 30/50
35015/35015 [=====] - 2s 43us/sample - loss: 40.3474 - mse: 3996.6465 - mae: 40.3474
Epoch 31/50
35015/35015 [=====] - 2s 43us/sample - loss: 40.3192 - mse: 3989.1357 - mae: 40.3192
Epoch 32/50
35015/35015 [=====] - 2s 43us/sample - loss: 40.2790 - mse: 3992.3813 - mae: 40.2790
Epoch 33/50
35015/35015 [=====] - 2s 43us/sample - loss: 40.2918 - mse: 3993.8240 - mae: 40.2918
Epoch 34/50

```

35015/35015 [=====] - 2s 44us/sample - loss: 40.2196 - mse: 3994.0249 - mae: 40.2196
Epoch 35/50
35015/35015 [=====] - 2s 43us/sample - loss: 40.2009 - mse: 3979.3872 - mae: 40.2009
Epoch 36/50
35015/35015 [=====] - 2s 43us/sample - loss: 40.0268 - mse: 3957.9846 - mae: 40.0269
Epoch 37/50
35015/35015 [=====] - 1s 43us/sample - loss: 39.8587 - mse: 3937.4265 - mae: 39.8587
Epoch 38/50
35015/35015 [=====] - 1s 43us/sample - loss: 40.0057 - mse: 3965.1650 - mae: 40.0057
Epoch 39/50
35015/35015 [=====] - 2s 43us/sample - loss: 39.7988 - mse: 3925.2485 - mae: 39.7988
Epoch 40/50
35015/35015 [=====] - 1s 43us/sample - loss: 39.7337 - mse: 3909.7881 - mae: 39.7337
Epoch 41/50
35015/35015 [=====] - 2s 44us/sample - loss: 39.6805 - mse: 3909.2004 - mae: 39.6805
Epoch 42/50
35015/35015 [=====] - 2s 44us/sample - loss: 39.5995 - mse: 3902.3271 - mae: 39.5995
Epoch 43/50
35015/35015 [=====] - 2s 43us/sample - loss: 39.5505 - mse: 3894.9324 - mae: 39.5505
Epoch 44/50
35015/35015 [=====] - 2s 43us/sample - loss: 39.5393 - mse: 3882.5322 - mae: 39.5393
Epoch 45/50
35015/35015 [=====] - 2s 43us/sample - loss: 39.4026 - mse: 3867.6694 - mae: 39.4026
Epoch 46/50
35015/35015 [=====] - 2s 44us/sample - loss: 39.4573 - mse: 3872.0000 - mae: 39.4573
Epoch 47/50
35015/35015 [=====] - 2s 44us/sample - loss: 39.3481 - mse: 3853.4172 - mae: 39.3481
Epoch 48/50
35015/35015 [=====] - 2s 43us/sample - loss: 39.1527 - mse: 3848.9690 - mae: 39.1527
Epoch 49/50
35015/35015 [=====] - 2s 43us/sample - loss: 39.3224 - mse: 3852.9966 - mae: 39.3224
Epoch 50/50
35015/35015 [=====] - 2s 43us/sample - loss: 39.2304 - mse: 3834.1753 - mae: 39.2304

```

```
plot_history(history)
```

```

# y_train = model.predict(X_train_gru)
# y_val = model.predict(X_val_gru)
y_train_all = model.predict(X_train_all_gru)

```


```
y_test = model.predict(x_test_gru)

# mae_train = mean_absolute_error(r_train_gru,y_train)
# mae_val = mean_absolute_error(r_val_gru,y_val)
mae_train_all = mean_absolute_error(r_train_all_gru,y_train_all)
mae_test = mean_absolute_error(r_test_gru,y_test)



# print(f"train      mae: {mae_train}")
# print(f"validation mae: {mae_val}")
print(f"train      mae: {mae_train_all}")
print(f"test mae: {mae_test}")


# Visualise first 1000 predictions for validation and test
# PlotResults(r_val_gru[:1000],y_val[:1000,0])
PlotResults(r_test_gru[:1000],y_test[:1000,0])
```





train mae: 38.36615753173828
test mae: 49.90842056274414





▼ Step 3

```
# current code for pollution-only, add other features for multi-feature model,
# remove indices as in commented code if you want to use all features
train_features = np.arange(14) # 16: all, 14: without snow and rain, 13: also leave pressure out, 12: also leave temperature out

CNN_WINDOW_SIZE = 2*168 # 2 weeks
ahead = 6
```

```
areaau = 0
```

```
# X_train_w,r_train_w = create_dataset_windowed(train_scaled[:,train_features], train_labels, window_size=CNN_WINDOW_SIZE)
# X_val_w,r_val_w = create_dataset_windowed(val_scaled[:,train_features], val_labels, window_size=CNN_WINDOW_SIZE)
X_train_all_w,r_train_all_w = create_dataset_windowed(train_all_scaled[:,train_features],train_all_labels, window_size=CNN_WINDOW_SIZ
X_test_w,r_test_w = create_dataset_windowed(test_scaled[:,train_features],test_labels, window_size=CNN_WINDOW_SIZE)

def conv_model(CNN_WINDOW_SIZE=24, learning_rate=0.01, hidden=[32,64,128,256]):
    model = Sequential()

    model.add(Conv1D(filters=hidden[0],
                     kernel_size=2,
                     padding='same',
                     kernel_initializer='he_uniform',
                     input_shape=(CNN_WINDOW_SIZE,len(train_features))))
    # model.add(BatchNormalization())
    model.add(Activation('relu'))
    model.add(MaxPooling1D(pool_size=2))

    for idx in range(1, len(hidden)):
        model.add(Conv1D(filters=hidden[idx],
                         kernel_size=2,
                         padding='same',
                         kernel_initializer='he_uniform'))
        # model.add(BatchNormalization())
        model.add(Activation('relu'))
        model.add(MaxPooling1D(pool_size=2))

    model.add(Flatten())
    model.add(Dense(1, kernel_initializer='glorot_uniform'))
    model.add(Activation('linear'))

    optim = tf.keras.optimizers.Adam(learning_rate=learning_rate)
    model.compile(loss='mae',
                  optimizer=optim,
                  # keep extra metrics: mse and mae without regularisation terms
                  metrics=['mse', 'mae'])
```

```
return model
```

```
batch_size = 24
epochs = 20
learning_rate = 0.0001

model = conv_model(CNN_WINDOW_SIZE=CNN_WINDOW_SIZE, learning_rate=learning_rate, hidden=[32,64,128,256])
model.summary()

# cp_path = '/content/gdrive/My Drive/Colab Notebooks/DL2020/GA3/checkpoints/step_3/cp.ckpt'
# cp_callback = tf.keras.callbacks.ModelCheckpoint(cp_path,
#
#                                     monitor='val_mae',
#                                     verbose=1,
#                                     save_best_only=True,
#                                     save_weights_only=True)

# history = model.fit(X_train_w, r_train_w,
#                     shuffle=True,
#                     batch_size=batch_size,
#                     epochs=epochs,
#                     validation_data=(X_val_w,r_val_w),
#                     callbacks=[cp_callback])
# model.load_weights(cp_path)

history = model.fit(X_train_all_w,
                    r_train_all_w,
                    shuffle=True,
                    batch_size=batch_size,
                    epochs=epochs)
```



Model: "sequential_3"

Layer (type)	Output Shape	Param #
=====		
conv1d_3 (Conv1D)	(None, 336, 32)	928
activation_10 (Activation)	(None, 336, 32)	0
max_pooling1d_3 (MaxPooling1D)	(None, 168, 32)	0
conv1d_4 (Conv1D)	(None, 168, 64)	4160
activation_11 (Activation)	(None, 168, 64)	0
max_pooling1d_4 (MaxPooling1D)	(None, 84, 64)	0
conv1d_5 (Conv1D)	(None, 84, 128)	16512
activation_12 (Activation)	(None, 84, 128)	0
max_pooling1d_5 (MaxPooling1D)	(None, 42, 128)	0
conv1d_6 (Conv1D)	(None, 42, 256)	65792
activation_13 (Activation)	(None, 42, 256)	0
max_pooling1d_6 (MaxPooling1D)	(None, 21, 256)	0
flatten_1 (Flatten)	(None, 5376)	0
dense_7 (Dense)	(None, 1)	5377
activation_14 (Activation)	(None, 1)	0

=====

Total params: 92,769
Trainable params: 92,769
Non-trainable params: 0

Train on 34699 samples

Epoch 1/20

34699/34699 [=====] - 8s 221us/sample - loss: 53.4678 - mse: 6566.6738 - mae: 53.4678

Epoch 2/20


```

34699/34699 [=====] - 6s 181us/sample - loss: 40.6294 - mse: 3859.4417 - mae: 40.6294
Epoch 3/20
34699/34699 [=====] - 6s 184us/sample - loss: 38.2617 - mse: 3574.7573 - mae: 38.2617
Epoch 4/20
34699/34699 [=====] - 6s 185us/sample - loss: 37.0370 - mse: 3434.5955 - mae: 37.0370
Epoch 5/20
34699/34699 [=====] - 6s 183us/sample - loss: 36.2029 - mse: 3347.6392 - mae: 36.2030
Epoch 6/20
34699/34699 [=====] - 6s 180us/sample - loss: 35.5974 - mse: 3270.3047 - mae: 35.5974
Epoch 7/20
34699/34699 [=====] - 7s 187us/sample - loss: 35.0903 - mse: 3217.6736 - mae: 35.0903
Epoch 8/20
34699/34699 [=====] - 7s 198us/sample - loss: 34.6493 - mse: 3166.6382 - mae: 34.6493
Epoch 9/20
34699/34699 [=====] - 7s 192us/sample - loss: 34.3211 - mse: 3121.6841 - mae: 34.3211
Epoch 10/20
34699/34699 [=====] - 6s 184us/sample - loss: 33.8776 - mse: 3079.3103 - mae: 33.8776
Epoch 11/20
34699/34699 [=====] - 6s 185us/sample - loss: 33.4903 - mse: 3039.8474 - mae: 33.4903
Epoch 12/20
34699/34699 [=====] - 6s 184us/sample - loss: 33.1742 - mse: 2996.2019 - mae: 33.1742
Epoch 13/20
34699/34699 [=====] - 6s 184us/sample - loss: 32.8745 - mse: 2960.5127 - mae: 32.8745
Epoch 14/20
34699/34699 [=====] - 6s 185us/sample - loss: 32.5650 - mse: 2924.0540 - mae: 32.5650
Epoch 15/20
34699/34699 [=====] - 6s 184us/sample - loss: 32.2604 - mse: 2883.9709 - mae: 32.2604
Epoch 16/20
34699/34699 [=====] - 6s 184us/sample - loss: 31.9605 - mse: 2843.4626 - mae: 31.9605
Epoch 17/20
34699/34699 [=====] - 7s 190us/sample - loss: 31.6633 - mse: 2809.0984 - mae: 31.6633
Epoch 18/20
34699/34699 [=====] - 6s 187us/sample - loss: 31.3906 - mse: 2778.4553 - mae: 31.3906
Epoch 19/20
34699/34699 [=====] - 6s 184us/sample - loss: 31.0827 - mse: 2738.3994 - mae: 31.0827
Epoch 20/20
34699/34699 [=====] - 6s 185us/sample - loss: 30.7971 - mse: 2704.5264 - mae: 30.7971

```

```
plot_history(history)
```


```
# y_train = model.predict(X_train_w)
# y_val = model.predict(X_val_w)
y_train_all = model.predict(X_train_all_w)
y_test = model.predict(X_test_w)

# mae_train = mean_absolute_error(r_train_w,y_train)
# mae_val = mean_absolute_error(r_val_w,y_val)
mae_train_all = mean_absolute_error(r_train_all_w,y_train_all)
mae_test = mean_absolute_error(r_test_w,y_test)



# print(f"train      mae: {mae_train}")
# print(f"validation mae: {mae_val}")
print(f"train all mae: {mae_train_all}")
print(f"test       mae: {mae_test}")


# Visualise first 1000 predictions for validation and test
# PlotResults(r_val_w[:1000],y_val[:1000,0])
PlotResults(r_test_w[:1000],y_test[:1000,0])
```





```
train all mae: 30.45968246459961
test      mae: 36.19908905029297
```





▼ Step 4

▼ Models

```
def conv_model(CNN_WINDOW_SIZE=24, learning_rate=0.01, hidden=[32,64,128,256], n_features=16):  
    model = Sequential()
```

```

model.add(Conv1D(filters=hidden[0],
                 kernel_size=2,
                 padding='same',
                 kernel_initializer='he_uniform',
                 input_shape=(CNN_WINDOW_SIZE,n_features)))
# model.add(BatchNormalization())
model.add(Activation('relu'))
model.add(MaxPooling1D(pool_size=2))

for idx in range(1, len(hidden)):
    model.add(Conv1D(filters=hidden[idx],
                     kernel_size=2,
                     padding='same',
                     kernel_initializer='he_uniform'))
    # model.add(BatchNormalization())
    model.add(Activation('relu'))
    model.add(MaxPooling1D(pool_size=2))

model.add(Flatten())
model.add(Dense(1, kernel_initializer='glorot_uniform'))
model.add(Activation('linear'))

optim = tf.keras.optimizers.Adam(learning_rate=learning_rate)
model.compile(loss='mae',
              optimizer=optim,
              # keep extra metrics: mse and mae without regularisation terms
              metrics=['mse', 'mae'])

return model

```

```

batch_size = 24
epochs = 20
learning_rate = 0.0001
CNN_WINDOW_SIZE = 2*168 # 2 weeks
ahead = 6

```

```

# 1
train_features = np.arange(16) # 16: all

# X_train_w, r_train_w = create_dataset_windowed(train_scaled[:,train_features], train_labels, window_size=CNN_WINDOW_SIZE)
# X_val_w, r_val_w = create_dataset_windowed(val_scaled[:,train_features], val_labels, window_size=CNN_WINDOW_SIZE)
X_train_all_w, r_train_all_w = create_dataset_windowed(train_all_scaled[:,train_features], train_all_labels, window_size=CNN_WINDOW_SIZ
X_test_w, r_test_w = create_dataset_windowed(test_scaled[:,train_features], test_labels, window_size=CNN_WINDOW_SIZE)

model = conv_model(CNN_WINDOW_SIZE=CNN_WINDOW_SIZE, learning_rate=learning_rate, hidden=[32,64,128,256], n_features=len(train_feature
# model.summary()

# cp_path = '/content/gdrive/My Drive/Colab Notebooks/DL2020/GA3/checkpoints/step_4/cp.ckpt'
# cp_callback = tf.keras.callbacks.ModelCheckpoint(cp_path,
#
#                                     monitor='val_mae',
#                                     verbose=1,
#                                     save_best_only=True,
#                                     save_weights_only=True)

# history = model.fit(X_train_w, r_train_w,
#                     shuffle=True,
#                     batch_size=batch_size,
#                     epochs=epochs,
#                     validation_data=(X_val_w, r_val_w),
#                     callbacks=[cp_callback])
# model.load_weights(cp_path)

history = model.fit(X_train_all_w, r_train_all_w, shuffle=True, batch_size=batch_size, epochs=epochs)

# y_train_1 = model.predict(X_train_w)
# y_val_1 = model.predict(X_val_w)
y_train_all_1 = model.predict(X_train_all_w)
y_test_1 = model.predict(X_test_w)

# mae_train = mean_absolute_error(r_train_w, y_train_1)
# mae_val = mean_absolute_error(r_val_w, y_val_1)
mae_train_all = mean_absolute_error(r_train_all_w, y_train_all_1)
mae_test = mean_absolute_error(r_test_w, y_test_1)

```

```
# print(f"train      mae: {mae_train}")
# print(f"validation mae: {mae_val}")
print(f"train all mae: {mae_train_all}")
print(f"test       mae: {mae_test}")
```



Train on 34699 samples

Epoch 1/20

34699/34699 [=====] - 7s 203us/sample - loss: 54.8534 - mse: 6857.3096 - mae: 54.8535

Epoch 2/20

34699/34699 [=====] - 6s 187us/sample - loss: 40.8519 - mse: 3880.9929 - mae: 40.8518

Epoch 3/20

34699/34699 [=====] - 7s 197us/sample - loss: 38.1572 - mse: 3535.9573 - mae: 38.1571

Epoch 4/20

34699/34699 [=====] - 7s 207us/sample - loss: 36.9534 - mse: 3386.5479 - mae: 36.9534

Epoch 5/20

34699/34699 [=====] - 7s 194us/sample - loss: 36.1111 - mse: 3287.3606 - mae: 36.1111

Epoch 6/20

34699/34699 [=====] - 7s 196us/sample - loss: 35.4299 - mse: 3210.5422 - mae: 35.4299

Epoch 7/20

34699/34699 [=====] - 7s 199us/sample - loss: 34.9252 - mse: 3159.7725 - mae: 34.9252

Epoch 8/20

34699/34699 [=====] - 7s 202us/sample - loss: 34.4660 - mse: 3102.1282 - mae: 34.4660

Epoch 9/20

34699/34699 [=====] - 7s 204us/sample - loss: 33.9536 - mse: 3044.1445 - mae: 33.9536

Epoch 10/20

34699/34699 [=====] - 7s 203us/sample - loss: 33.6187 - mse: 3009.4475 - mae: 33.6187

Epoch 11/20

34699/34699 [=====] - 7s 205us/sample - loss: 33.1836 - mse: 2955.1345 - mae: 33.1835

Epoch 12/20

34699/34699 [=====] - 7s 203us/sample - loss: 32.8078 - mse: 2910.6738 - mae: 32.8078

Epoch 13/20

34699/34699 [=====] - 7s 200us/sample - loss: 32.4859 - mse: 2867.5222 - mae: 32.4860

Epoch 14/20

34699/34699 [=====] - 7s 194us/sample - loss: 32.1952 - mse: 2826.7388 - mae: 32.1952

Epoch 15/20

34699/34699 [=====] - 7s 196us/sample - loss: 31.7864 - mse: 2783.0283 - mae: 31.7864

Epoch 16/20

34699/34699 [=====] - 7s 196us/sample - loss: 31.4701 - mse: 2743.9824 - mae: 31.4700

Epoch 17/20

34699/34699 [=====] - 7s 195us/sample - loss: 31.1860 - mse: 2707.5129 - mae: 31.1860

Epoch 18/20

34699/34699 [=====] - 7s 194us/sample - loss: 30.8736 - mse: 2669.2893 - mae: 30.8735

Epoch 19/20

34699/34699 [=====] - 7s 196us/sample - loss: 30.5661 - mse: 2630.2693 - mae: 30.5661

Epoch 20/20

34699/34699 [=====] - 7s 194us/sample - loss: 30.2766 - mse: 2597.5564 - mae: 30.2766

train all mae: 30.693071365356445


```
test      mae: 36.721683502197266
```

```
# 2
train_features = np.arange(14) # 14: without snow and rain

# X_train_w, r_train_w = create_dataset_windowed(train_scaled[:,train_features], train_labels, window_size=CNN_WINDOW_SIZE)
# X_val_w, r_val_w = create_dataset_windowed(val_scaled[:,train_features], val_labels, window_size=CNN_WINDOW_SIZE)
X_train_all_w, r_train_all_w = create_dataset_windowed(train_all_scaled[:,train_features], train_all_labels, window_size=CNN_WINDOW_SIZE)
X_test_w, r_test_w = create_dataset_windowed(test_scaled[:,train_features], test_labels, window_size=CNN_WINDOW_SIZE)

model = conv_model(CNN_WINDOW_SIZE=CNN_WINDOW_SIZE, learning_rate=learning_rate, hidden=[32,64,128,256], n_features=len(train_features))
# model.summary()

# cp_path = '/content/gdrive/My Drive/Colab Notebooks/DL2020/GA3/checkpoints/step_4/cp.ckpt'
# cp_callback = tf.keras.callbacks.ModelCheckpoint(cp_path,
#
#                                     monitor='val_mae',
#                                     verbose=1,
#                                     save_best_only=True,
#                                     save_weights_only=True)

# history = model.fit(X_train_w, r_train_w,
#                     shuffle=True,
#                     batch_size=batch_size,
#                     epochs=epochs,
#                     validation_data=(X_val_w, r_val_w),
#                     callbacks=[cp_callback])
# model.load_weights(cp_path)

history = model.fit(X_train_all_w, r_train_all_w, shuffle=True, batch_size=batch_size, epochs=epochs)

# y_train_2 = model.predict(X_train_w)
# y_val_2 = model.predict(X_val_w)
y_train_all_2 = model.predict(X_train_all_w)
y_test_2 = model.predict(X_test_w)

# mae_train = mean_absolute_error(r_train_w, y_train_2)
# mae_val = mean_absolute_error(r_val_w, y_val_2)
```

```
# mae_val = mean_absolute_error(r_val_w,y_val_2)
mae_train_all = mean_absolute_error(r_train_all_w,y_train_all_2)
mae_test = mean_absolute_error(r_test_w,y_test_2)

# print(f"train      mae: {mae_train}")
# print(f"validation mae: {mae_val}")
print(f"train all mae: {mae_train_all}")
print(f"test       mae: {mae_test}")
```



Train on 34699 samples

Epoch 1/20

34699/34699 [=====] - 7s 212us/sample - loss: 52.8202 - mse: 6273.0283 - mae: 52.8202

Epoch 2/20

34699/34699 [=====] - 7s 197us/sample - loss: 40.0717 - mse: 3761.2554 - mae: 40.0717

Epoch 3/20

34699/34699 [=====] - 7s 199us/sample - loss: 37.8851 - mse: 3519.9519 - mae: 37.8851

Epoch 4/20

34699/34699 [=====] - 7s 196us/sample - loss: 36.8386 - mse: 3413.6392 - mae: 36.8386

Epoch 5/20

34699/34699 [=====] - 7s 194us/sample - loss: 36.1003 - mse: 3347.7534 - mae: 36.1004

Epoch 6/20

34699/34699 [=====] - 7s 193us/sample - loss: 35.4671 - mse: 3277.6780 - mae: 35.4671

Epoch 7/20

34699/34699 [=====] - 7s 195us/sample - loss: 34.9650 - mse: 3222.1902 - mae: 34.9650

Epoch 8/20

34699/34699 [=====] - 7s 197us/sample - loss: 34.4782 - mse: 3171.0457 - mae: 34.4783

Epoch 9/20

34699/34699 [=====] - 7s 197us/sample - loss: 34.0872 - mse: 3125.9915 - mae: 34.0872

Epoch 10/20

34699/34699 [=====] - 7s 194us/sample - loss: 33.6866 - mse: 3076.1357 - mae: 33.6866

Epoch 11/20

34699/34699 [=====] - 7s 194us/sample - loss: 33.3307 - mse: 3041.9434 - mae: 33.3308

Epoch 12/20

34699/34699 [=====] - 7s 195us/sample - loss: 32.9633 - mse: 3003.2407 - mae: 32.9634

Epoch 13/20

34699/34699 [=====] - 7s 203us/sample - loss: 32.5760 - mse: 2952.4656 - mae: 32.5761

Epoch 14/20

34699/34699 [=====] - 7s 195us/sample - loss: 32.2687 - mse: 2924.3916 - mae: 32.2687

Epoch 15/20

34699/34699 [=====] - 7s 196us/sample - loss: 31.9347 - mse: 2884.2375 - mae: 31.9346

Epoch 16/20

34699/34699 [=====] - 7s 194us/sample - loss: 31.6328 - mse: 2843.5540 - mae: 31.6328

Epoch 17/20

34699/34699 [=====] - 7s 194us/sample - loss: 31.3857 - mse: 2811.0469 - mae: 31.3857

Epoch 18/20

34699/34699 [=====] - 7s 195us/sample - loss: 31.0480 - mse: 2773.1462 - mae: 31.0479

Epoch 19/20

34699/34699 [=====] - 7s 192us/sample - loss: 30.7533 - mse: 2741.5356 - mae: 30.7533

Epoch 20/20

34699/34699 [=====] - 7s 193us/sample - loss: 30.4540 - mse: 2708.9795 - mae: 30.4540

train all mae: 30.0151309967041

```
test      mae: 36.05653381347656
```

```
# 3
train_features = np.arange(13) # 13: also leave pressure out

# X_train_w, r_train_w = create_dataset_windowed(train_scaled[:,train_features], train_labels, window_size=CNN_WINDOW_SIZE)
# X_val_w, r_val_w = create_dataset_windowed(val_scaled[:,train_features], val_labels, window_size=CNN_WINDOW_SIZE)
X_train_all_w, r_train_all_w = create_dataset_windowed(train_all_scaled[:,train_features], train_all_labels, window_size=CNN_WINDOW_SIZE)
X_test_w, r_test_w = create_dataset_windowed(test_scaled[:,train_features], test_labels, window_size=CNN_WINDOW_SIZE)

model = conv_model(CNN_WINDOW_SIZE=CNN_WINDOW_SIZE, learning_rate=learning_rate, hidden=[32,64,128,256], n_features=len(train_features))
# model.summary()

# cp_path = '/content/gdrive/My Drive/Colab Notebooks/DL2020/GA3/checkpoints/step_4/cp.ckpt'
# cp_callback = tf.keras.callbacks.ModelCheckpoint(cp_path,
#
#                                     monitor='val_mae',
#                                     verbose=1,
#                                     save_best_only=True,
#                                     save_weights_only=True)

# history = model.fit(X_train_w, r_train_w,
#                     shuffle=True,
#                     batch_size=batch_size,
#                     epochs=epochs,
#                     validation_data=(X_val_w, r_val_w),
#                     callbacks=[cp_callback])
# model.load_weights(cp_path)

history = model.fit(X_train_all_w, r_train_all_w, shuffle=True, batch_size=batch_size, epochs=epochs)

# y_train_3 = model.predict(X_train_w)
# y_val_3 = model.predict(X_val_w)
y_train_all_3 = model.predict(X_train_all_w)
y_test_3 = model.predict(X_test_w)

# mae_train = mean_absolute_error(r_train_w, y_train_3)
# mae_val = mean_absolute_error(r_val_w, y_val_3)
```

```
# mae_val = mean_absolute_error(r_val_w,y_val_3)
mae_train_all = mean_absolute_error(r_train_all_w,y_train_all_3)
mae_test = mean_absolute_error(r_test_w,y_test_3)

# print(f"train      mae: {mae_train}")
# print(f"validation mae: {mae_val}")
print(f"train all mae: {mae_train_all}")
print(f"test       mae: {mae_test}")
```



Train on 34699 samples

Epoch 1/20

34699/34699 [=====] - 7s 210us/sample - loss: 54.0307 - mse: 6617.4917 - mae: 54.0307

Epoch 2/20

34699/34699 [=====] - 7s 190us/sample - loss: 40.2426 - mse: 3807.7834 - mae: 40.2426

Epoch 3/20

34699/34699 [=====] - 7s 191us/sample - loss: 37.9463 - mse: 3515.8484 - mae: 37.9462

Epoch 4/20

34699/34699 [=====] - 7s 191us/sample - loss: 36.8989 - mse: 3408.2424 - mae: 36.8989

Epoch 5/20

34699/34699 [=====] - 7s 191us/sample - loss: 36.1279 - mse: 3325.7302 - mae: 36.1279

Epoch 6/20

34699/34699 [=====] - 7s 193us/sample - loss: 35.6114 - mse: 3265.9702 - mae: 35.6114

Epoch 7/20

34699/34699 [=====] - 7s 196us/sample - loss: 35.1045 - mse: 3212.6860 - mae: 35.1044

Epoch 8/20

34699/34699 [=====] - 7s 191us/sample - loss: 34.7127 - mse: 3173.8320 - mae: 34.7126

Epoch 9/20

34699/34699 [=====] - 7s 190us/sample - loss: 34.2677 - mse: 3125.3076 - mae: 34.2677

Epoch 10/20

34699/34699 [=====] - 7s 189us/sample - loss: 33.9107 - mse: 3088.5574 - mae: 33.9106

Epoch 11/20

34699/34699 [=====] - 7s 190us/sample - loss: 33.5652 - mse: 3050.0444 - mae: 33.5652

Epoch 12/20

34699/34699 [=====] - 7s 191us/sample - loss: 33.1941 - mse: 3000.4568 - mae: 33.1941

Epoch 13/20

34699/34699 [=====] - 7s 193us/sample - loss: 32.8671 - mse: 2969.6519 - mae: 32.8671

Epoch 14/20

34699/34699 [=====] - 7s 192us/sample - loss: 32.5818 - mse: 2933.4343 - mae: 32.5818

Epoch 15/20

34699/34699 [=====] - 7s 192us/sample - loss: 32.2261 - mse: 2899.2178 - mae: 32.2261

Epoch 16/20

34699/34699 [=====] - 7s 193us/sample - loss: 31.9384 - mse: 2861.4075 - mae: 31.9384

Epoch 17/20

34699/34699 [=====] - 7s 192us/sample - loss: 31.6317 - mse: 2831.4285 - mae: 31.6318

Epoch 18/20

34699/34699 [=====] - 7s 192us/sample - loss: 31.3363 - mse: 2790.1589 - mae: 31.3363

Epoch 19/20

34699/34699 [=====] - 7s 193us/sample - loss: 31.0215 - mse: 2754.9412 - mae: 31.0215

Epoch 20/20

34699/34699 [=====] - 7s 191us/sample - loss: 30.7538 - mse: 2717.4949 - mae: 30.7538

train all mae: 30.32161521911621

```
test      mae: 35.549591064453125
```

```
# 4
train_features = np.arange(12) # 12: also leave temperature out

# X_train_w, r_train_w = create_dataset_windowed(train_scaled[:,train_features], train_labels, window_size=CNN_WINDOW_SIZE)
# X_val_w, r_val_w = create_dataset_windowed(val_scaled[:,train_features], val_labels, window_size=CNN_WINDOW_SIZE)
X_train_all_w, r_train_all_w = create_dataset_windowed(train_all_scaled[:,train_features], train_all_labels, window_size=CNN_WINDOW_SIZE)
X_test_w, r_test_w = create_dataset_windowed(test_scaled[:,train_features], test_labels, window_size=CNN_WINDOW_SIZE)

model = conv_model(CNN_WINDOW_SIZE=CNN_WINDOW_SIZE, learning_rate=learning_rate, hidden=[32,64,128,256], n_features=len(train_features))
# model.summary()

# cp_path = '/content/gdrive/My Drive/Colab Notebooks/DL2020/GA3/checkpoints/step_4/cp.ckpt'
# cp_callback = tf.keras.callbacks.ModelCheckpoint(cp_path,
#
#                                     monitor='val_mae',
#                                     verbose=1,
#                                     save_best_only=True,
#                                     save_weights_only=True)

# history = model.fit(X_train_w, r_train_w,
#                     shuffle=True,
#                     batch_size=batch_size,
#                     epochs=epochs,
#                     validation_data=(X_val_w, r_val_w),
#                     callbacks=[cp_callback])
# model.load_weights(cp_path)

history = model.fit(X_train_all_w, r_train_all_w, shuffle=True, batch_size=batch_size, epochs=epochs)

# y_train_4 = model.predict(X_train_w)
# y_val_4 = model.predict(X_val_w)
y_train_all_4 = model.predict(X_train_all_w)
y_test_4 = model.predict(X_test_w)

# mae_train = mean_absolute_error(r_train_w, y_train_4)
# mae_val = mean_absolute_error(r_val_w, y_val_4)
```

```
# mae_val = mean_absolute_error(r_val_w,y_val_4)
mae_train_all = mean_absolute_error(r_train_all_w,y_train_all_4)
mae_test = mean_absolute_error(r_test_w,y_test_4)

# print(f"train      mae: {mae_train}")
# print(f"validation mae: {mae_val}")
print(f"train all mae: {mae_train_all}")
print(f"test       mae: {mae_test}")
```



Train on 34699 samples

Epoch 1/20

34699/34699 [=====] - 7s 207us/sample - loss: 52.7946 - mse: 6387.0820 - mae: 52.7945

Epoch 2/20

34699/34699 [=====] - 7s 190us/sample - loss: 40.1069 - mse: 3812.5237 - mae: 40.1069

Epoch 3/20

34699/34699 [=====] - 7s 190us/sample - loss: 37.9810 - mse: 3552.3857 - mae: 37.9810

Epoch 4/20

34699/34699 [=====] - 7s 191us/sample - loss: 37.0116 - mse: 3439.1738 - mae: 37.0116

Epoch 5/20

34699/34699 [=====] - 7s 189us/sample - loss: 36.3666 - mse: 3364.1562 - mae: 36.3666

Epoch 6/20

34699/34699 [=====] - 7s 190us/sample - loss: 35.8381 - mse: 3306.1174 - mae: 35.8381

Epoch 7/20

34699/34699 [=====] - 7s 194us/sample - loss: 35.4023 - mse: 3251.7136 - mae: 35.4023

Epoch 8/20

34699/34699 [=====] - 7s 190us/sample - loss: 34.9676 - mse: 3201.4102 - mae: 34.9676

Epoch 9/20

34699/34699 [=====] - 7s 194us/sample - loss: 34.6092 - mse: 3158.4285 - mae: 34.6092

Epoch 10/20

34699/34699 [=====] - 7s 198us/sample - loss: 34.2632 - mse: 3128.3069 - mae: 34.2632

Epoch 11/20

34699/34699 [=====] - 7s 192us/sample - loss: 33.9271 - mse: 3084.5957 - mae: 33.9271

Epoch 12/20

34699/34699 [=====] - 7s 191us/sample - loss: 33.5960 - mse: 3039.9888 - mae: 33.5959

Epoch 13/20

34699/34699 [=====] - 7s 191us/sample - loss: 33.2688 - mse: 3004.7747 - mae: 33.2688

Epoch 14/20

34699/34699 [=====] - 6s 186us/sample - loss: 32.9851 - mse: 2965.5227 - mae: 32.9851

Epoch 15/20

34699/34699 [=====] - 7s 190us/sample - loss: 32.6499 - mse: 2937.2498 - mae: 32.6499

Epoch 16/20

34699/34699 [=====] - 7s 192us/sample - loss: 32.3263 - mse: 2891.3926 - mae: 32.3264

Epoch 17/20

34699/34699 [=====] - 7s 188us/sample - loss: 31.9796 - mse: 2857.7935 - mae: 31.9796

Epoch 18/20

34699/34699 [=====] - 7s 197us/sample - loss: 31.7465 - mse: 2819.9917 - mae: 31.7465

Epoch 19/20

34699/34699 [=====] - 7s 196us/sample - loss: 31.4289 - mse: 2783.5835 - mae: 31.4289

Epoch 20/20

34699/34699 [=====] - 7s 188us/sample - loss: 31.1717 - mse: 2758.3677 - mae: 31.1717

train all mae: 30.537208557128906

```
test      mae: 36.466339111328125
```

```
# 5
train_features = [0,] # only pollution

# X_train_w, r_train_w = create_dataset_windowed(train_scaled[:,train_features], train_labels, window_size=CNN_WINDOW_SIZE)
# X_val_w, r_val_w = create_dataset_windowed(val_scaled[:,train_features], val_labels, window_size=CNN_WINDOW_SIZE)
X_train_all_w, r_train_all_w = create_dataset_windowed(train_all_scaled[:,train_features], train_all_labels, window_size=CNN_WINDOW_SIZE)
X_test_w, r_test_w = create_dataset_windowed(test_scaled[:,train_features], test_labels, window_size=CNN_WINDOW_SIZE)

model = conv_model(CNN_WINDOW_SIZE=CNN_WINDOW_SIZE, learning_rate=learning_rate, hidden=[32,64,128,256], n_features=len(train_features))
# model.summary()

# cp_path = '/content/gdrive/My Drive/Colab Notebooks/DL2020/GA3/checkpoints/step_4/cp.ckpt'
# cp_callback = tf.keras.callbacks.ModelCheckpoint(cp_path,
#
#                                     monitor='val_mae',
#                                     verbose=1,
#                                     save_best_only=True,
#                                     save_weights_only=True)

# history = model.fit(X_train_w, r_train_w,
#                     shuffle=True,
#                     batch_size=batch_size,
#                     epochs=epochs,
#                     validation_data=(X_val_w, r_val_w),
#                     callbacks=[cp_callback])
# model.load_weights(cp_path)

history = model.fit(X_train_all_w, r_train_all_w, shuffle=True, batch_size=batch_size, epochs=epochs)

# y_train_5 = model.predict(X_train_w)
# y_val_5 = model.predict(X_val_w)
y_train_all_5 = model.predict(X_train_all_w)
y_test_5 = model.predict(X_test_w)

# mae_train = mean_absolute_error(r_train_w, y_train_5)
# mae_val = mean_absolute_error(r_val_w, y_val_5)
```

```
# mae_val = mean_absolute_error(r_val_w,y_val_5)
mae_train_all = mean_absolute_error(r_train_all_w,y_train_all_5)
mae_test = mean_absolute_error(r_test_w,y_test_5)

# print(f"train      mae: {mae_train}")
# print(f"validation mae: {mae_val}")
print(f"train all mae: {mae_train_all}")
print(f"test       mae: {mae_test}")
```



Train on 34699 samples

Epoch 1/20

34699/34699 [=====] - 7s 190us/sample - loss: 50.5201 - mse: 5662.3936 - mae: 50.5200

Epoch 2/20

34699/34699 [=====] - 6s 172us/sample - loss: 41.0793 - mse: 3996.4307 - mae: 41.0793

Epoch 3/20

34699/34699 [=====] - 6s 170us/sample - loss: 39.6769 - mse: 3839.4705 - mae: 39.6769

Epoch 4/20

34699/34699 [=====] - 6s 173us/sample - loss: 38.8301 - mse: 3750.6023 - mae: 38.8301

Epoch 5/20

34699/34699 [=====] - 6s 172us/sample - loss: 38.3320 - mse: 3694.9348 - mae: 38.3320

Epoch 6/20

34699/34699 [=====] - 6s 169us/sample - loss: 38.0554 - mse: 3672.5378 - mae: 38.0554

Epoch 7/20

34699/34699 [=====] - 6s 172us/sample - loss: 37.7645 - mse: 3641.6311 - mae: 37.7645

Epoch 8/20

34699/34699 [=====] - 6s 170us/sample - loss: 37.5976 - mse: 3621.0688 - mae: 37.5977

Epoch 9/20

34699/34699 [=====] - 6s 170us/sample - loss: 37.4052 - mse: 3604.5259 - mae: 37.4052

Epoch 10/20

34699/34699 [=====] - 6s 172us/sample - loss: 37.3063 - mse: 3583.9077 - mae: 37.3063

Epoch 11/20

34699/34699 [=====] - 6s 171us/sample - loss: 37.0894 - mse: 3572.0010 - mae: 37.0894

Epoch 12/20

34699/34699 [=====] - 6s 175us/sample - loss: 36.9698 - mse: 3555.3555 - mae: 36.9698

Epoch 13/20

34699/34699 [=====] - 6s 176us/sample - loss: 36.8417 - mse: 3530.8772 - mae: 36.8417

Epoch 14/20

34699/34699 [=====] - 6s 175us/sample - loss: 36.7225 - mse: 3527.4028 - mae: 36.7225

Epoch 15/20

34699/34699 [=====] - 6s 176us/sample - loss: 36.5896 - mse: 3510.2017 - mae: 36.5896

Epoch 16/20

34699/34699 [=====] - 6s 173us/sample - loss: 36.4849 - mse: 3485.8623 - mae: 36.4849

Epoch 17/20

34699/34699 [=====] - 6s 173us/sample - loss: 36.3171 - mse: 3470.1851 - mae: 36.3171

Epoch 18/20

34699/34699 [=====] - 6s 173us/sample - loss: 36.2053 - mse: 3465.5117 - mae: 36.2053

Epoch 19/20

34699/34699 [=====] - 6s 177us/sample - loss: 36.1076 - mse: 3448.4753 - mae: 36.1077

Epoch 20/20

34699/34699 [=====] - 6s 179us/sample - loss: 36.0124 - mse: 3434.7559 - mae: 36.0124

train all mae: 35.67976379394531

```
test      mae: 37.383750915527344
```

► Average

↳ 1 cell hidden

▼ Weighted average

```
# create new dataset
# X_train_t,r_train_t = create_dataset_top((y_train_1, y_train_2, y_train_3, y_train_4, y_train_5), train_labels, window_size=CNN_WIN
# X_val_t,r_val_t = create_dataset_top((y_val_1, y_val_2, y_val_3, y_val_4, y_val_5), val_labels, window_size=CNN_WINDOW_SIZE)
X_train_all_t,r_train_all_t = create_dataset_top((y_train_all_1, y_train_all_2, y_train_all_3, y_train_all_4, y_train_all_5), train_a
X_test_t,r_test_t = create_dataset_top((y_test_1, y_test_2, y_test_3, y_test_4, y_test_5), test_labels, window_size=CNN_WINDOW_SIZE)

def top_model(n_models=5, learning_rate=0.01):
    model = Sequential(Dense(1, activation='linear', kernel_initializer=tf.constant_initializer(value=1/n_models), kernel_constraint=tf

    optim = tf.keras.optimizers.Adam(learning_rate=learning_rate)
    model.compile(loss='mae',
                  optimizer=optim,
                  # keep extra metrics: mse and mae without regularisation terms
                  metrics=['mse', 'mae'])

    return model

batch_size = 32
epochs = 20
learning_rate = 0.001

model = top_model(n_models=5, learning_rate=learning_rate)
model.summary()

# cp_path = '/content/gdrive/My Drive/Colab Notebooks/DL2020/GA3/checkpoints/step_4/cp.ckpt'
```

```
# cp_callback = tf.keras.callbacks.ModelCheckpoint(cp_path,
#
#                                     monitor='val_mae',
#                                     verbose=1,
#                                     save_best_only=True,
#                                     save_weights_only=True)

# history = model.fit(X_train_t, r_train_t,
#                     shuffle=True,
#                     batch_size=batch_size,
#                     epochs=epochs,
#                     validation_data=(X_val_t,r_val_t),
#                     callbacks=[cp_callback])
# model.load_weights(cp_path)

history = model.fit(X_train_all_t, r_train_all_t, shuffle=True, batch_size=batch_size, epochs=epochs)
```



Model: "sequential_9"

Layer (type)	Output Shape	Param #
dense_13 (Dense)	(None, 1)	5

Total params: 5

Trainable params: 5

Non-trainable params: 0

Train on 34699 samples

Epoch 1/20

34699/34699 [=====] - 3s 89us/sample - loss: 29.5108 - mse: 2583.5007 - mae: 29.5108

Epoch 2/20

34699/34699 [=====] - 3s 85us/sample - loss: 29.2769 - mse: 2557.5396 - mae: 29.2769

Epoch 3/20

34699/34699 [=====] - 3s 82us/sample - loss: 29.1851 - mse: 2542.3557 - mae: 29.1851

Epoch 4/20

34699/34699 [=====] - 3s 83us/sample - loss: 29.1535 - mse: 2533.5254 - mae: 29.1535

Epoch 5/20

34699/34699 [=====] - 3s 83us/sample - loss: 29.1325 - mse: 2534.1399 - mae: 29.1325

Epoch 6/20

34699/34699 [=====] - 3s 83us/sample - loss: 29.1342 - mse: 2531.3281 - mae: 29.1342

Epoch 7/20

34699/34699 [=====] - 3s 84us/sample - loss: 29.1299 - mse: 2531.9302 - mae: 29.1299

Epoch 8/20

34699/34699 [=====] - 3s 82us/sample - loss: 29.1212 - mse: 2530.6082 - mae: 29.1212

Epoch 9/20

34699/34699 [=====] - 3s 82us/sample - loss: 29.1216 - mse: 2528.7749 - mae: 29.1217

Epoch 10/20

34699/34699 [=====] - 3s 83us/sample - loss: 29.1191 - mse: 2528.1057 - mae: 29.1191

Epoch 11/20

34699/34699 [=====] - 3s 83us/sample - loss: 29.1208 - mse: 2528.4241 - mae: 29.1208

Epoch 12/20

34699/34699 [=====] - 3s 83us/sample - loss: 29.1156 - mse: 2528.0950 - mae: 29.1156

Epoch 13/20

34699/34699 [=====] - 3s 82us/sample - loss: 29.1133 - mse: 2528.0969 - mae: 29.1133

Epoch 14/20

34699/34699 [=====] - 3s 83us/sample - loss: 29.1043 - mse: 2527.8591 - mae: 29.1043

Epoch 15/20

34699/34699 [=====] - 3s 82us/sample - loss: 29.1180 - mse: 2528.7607 - mae: 29.1180

Epoch 16/20

```

34699/34699 [=====] - 3s 81us/sample - loss: 29.1141 - mse: 2526.6499 - mae: 29.1141
Epoch 17/20
34699/34699 [=====] - 3s 83us/sample - loss: 29.1115 - mse: 2524.7542 - mae: 29.1115
Epoch 18/20
34699/34699 [=====] - 3s 82us/sample - loss: 29.1182 - mse: 2528.4539 - mae: 29.1183
Epoch 19/20
34699/34699 [=====] - 3s 83us/sample - loss: 29.1034 - mse: 2525.4360 - mae: 29.1034
Epoch 20/20
34699/34699 [=====] - 3s 83us/sample - loss: 29.1080 - mse: 2525.6501 - mae: 29.1080

```

```

plot_history(history)

# y_train = model.predict(X_train_t)
# y_val = model.predict(X_val_t)
y_train_all = model.predict(X_train_all_t)
y_test = model.predict(X_test_t)


# mae_train = mean_absolute_error(r_train_t,y_train)
# mae_val = mean_absolute_error(r_val_t,y_val)
mae_train_all = mean_absolute_error(r_train_all_t,y_train_all)
mae_test = mean_absolute_error(r_test_t,y_test)

# print(f"train      mae: {mae_train}")
# print(f"validation mae: {mae_val}")
print(f"train all mae: {mae_train_all}")
print(f"test      mae: {mae_test}")



# Visualise first 1000 predictions for test
# PlotResults(r_val_t[:1000],y_val[:1000,0])
PlotResults(r_test_t[:1000],y_test[:1000,0])


```





train all mae: 29.08261489868164
test mae: 35.12261962890625





```
# get weights of each model  
model.layers[0].get_weights()
```

```
↳ array([[ 0.31255028],  
         [ 0.34277973],  
         [ 0.34602535],  
         [ 0.06957645],  
         [-0.08152214]], dtype=float32)]
```

